

# Notes and Drafts for the Paper

Zuo Xiang

September 17, 2018

## 1 NFV Infrastructure

Motivation: Given an overview of the NFV infrastructure and then indicate which parts introduce latencies.

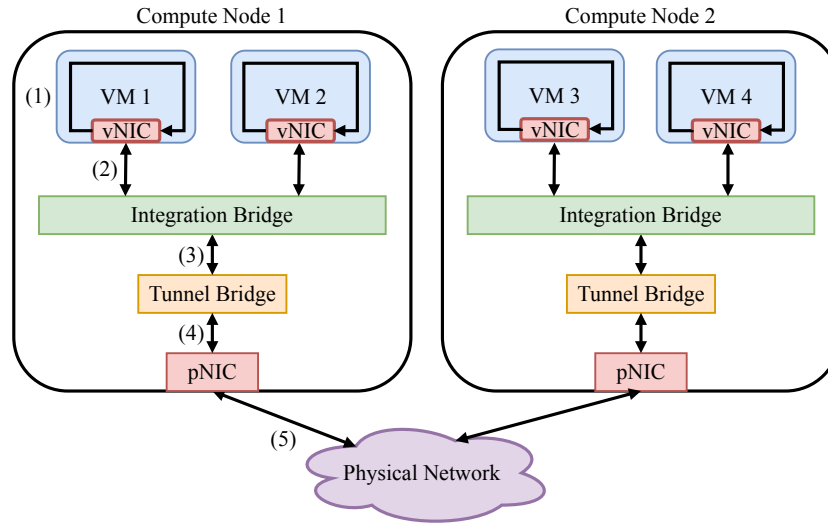


Figure 1: Common components of a compute node in a NFV infrastructure

The Fig.1 shows components of a compute node (can be seen as a physical elementary processing unit in a cloud environment). Multiple Virtual Instances (VIs, can be virtual machine or container, but currently VMs are still the most deployed and default option.) can be created and connected on different compute nodes to provide services. In order to provide a virtual network overlay on top of the underlying physical provider network,

additional bridges (most used are virtual bridge implementations e.g. Open VSwitch) are used to provide network connectivity among VIs. For example, VI1 and VI4 can be configured in the same virtual LAN network (in the same broadcast domain) even they are actually on different physical nodes which can be in the different physical LANs. (Or to say L3 accessible through routers). If VI1 send a packet to VI4, the packet needs to go through all components (1-2-3-4-5-6-7-8) where additional latencies are incurred. Since our implementation uses virtual machines, I will use VM instead of VI in following parts.

Introduction of the path and incurred latencies:

- (1) and (8) - From VI to the integration bridge (br-int) and inverse:

**This is the latency part that we have significantly reduced with proposed approach.**

Packets are transmitted to the VM, **processed by the VM** and sent back to the br-int. According to some survey paper and industry reports, with the performance improvements of virtual bridges, this part becomes now a bottleneck in terms of latency. (In other words, processing packets in the VM costs too much time.) How our approach reduce the latency will be described in 2.

- (2)-(3) and (6)-(7) - From virtual network overlay to physical network and inverse:

Latency here is mainly reduced by using **OVS-DPDK** (Open vSwitch with DPDK as Datapath) [OVS-DPDK Homepage](#). We have deployed OVS-DPDK on our testbed but without any modification or optimization.

- (4)-(5) - Physical network transmission: The latency here depends on the physical network technologies which are not optimized in this paper. All (three) physical compute nodes are connected with Gigabit Ethernet switches (in the same LAN), all layer 3 services (DHCP, L3 routing etc.) are provided by an additional network node in the same LAN. Latency here is **negligible** because of good hardware, small setup, small traffic congestion.

## 2 Approaches

Comment: Names of following approaches is to be discussed.

Before explaining different approaches, the logical service loop is illustrated in Fig 2. It shows how each packet is received by the cloud, processed by a chain of networking services (this is the concept of service function chain), transmitted to the requested server, processed by the server and leave the cloud. The egress SFC is mostly optional.

The ingress and egress points are endpoints or instances that are exposed to the external network (outside of the cloud network which is commonly a dedicated private network.). For e.g. security considerations, not all internal components are exposed to the external network. Therefore, ingress and egress points are required for clients in the external network to have access to the cloud services. Here the components of the ingress service chain shows the common logical function units of a service chain, namely an ordered set of virtualized network functions (VNFs). Some of these functions can be performed parallel and some of them required to be piped and work as a pipeline.

Because the latency before the packet arriving at the ingress point and after it leaving the egress point is fully depends on the external network setup(e.g. cable or fiber are used. it's out of scope of our optimizations). We have measured the latency between ingress and egress points with **only one** VNF and a server in the middle. This is the RTT value plotted in the Fig 3. To make things simple(or independent from the service application), the server here is just a bounce-server: receive a packet and send this packet out without doing anything. And the VNFs in the chain perform the elementary forwarding function (Can also perform network coding, XOR and simple encryption). Based on this loop, three approaches are explained as following.

### 2.0.1 Direct Connection or Dummy Forwarding

As plotted in the Fig 2, this means directly connect(shorten) the ingress and egress points. With this approach, the cloud does nothing on packets. This can be used as the baseline to evaluate the overhead introduced by following approaches. This can be also used as an indicator for our basic setup. It shows that our OpenStack cluster introduce already about 100 microseconds (us).

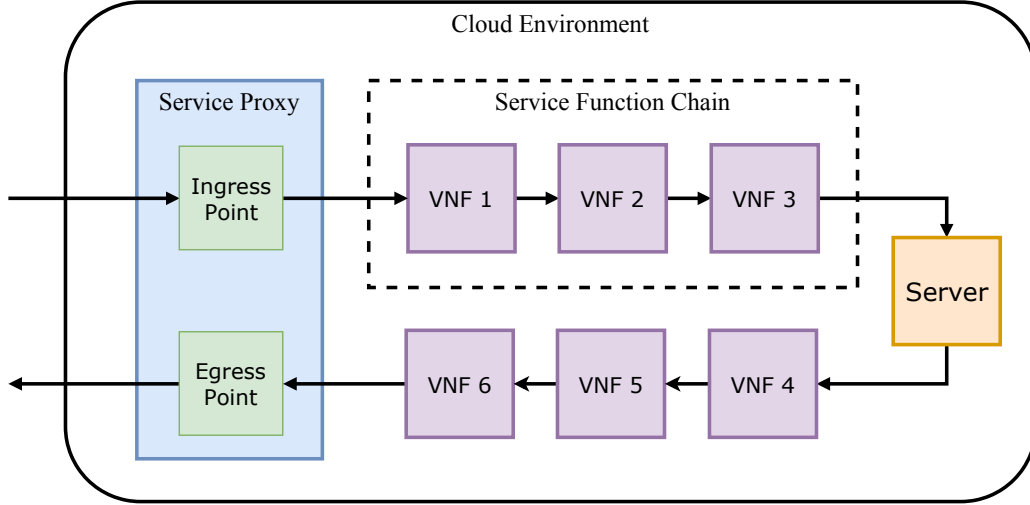


Figure 2: Service Loop inside Cloud

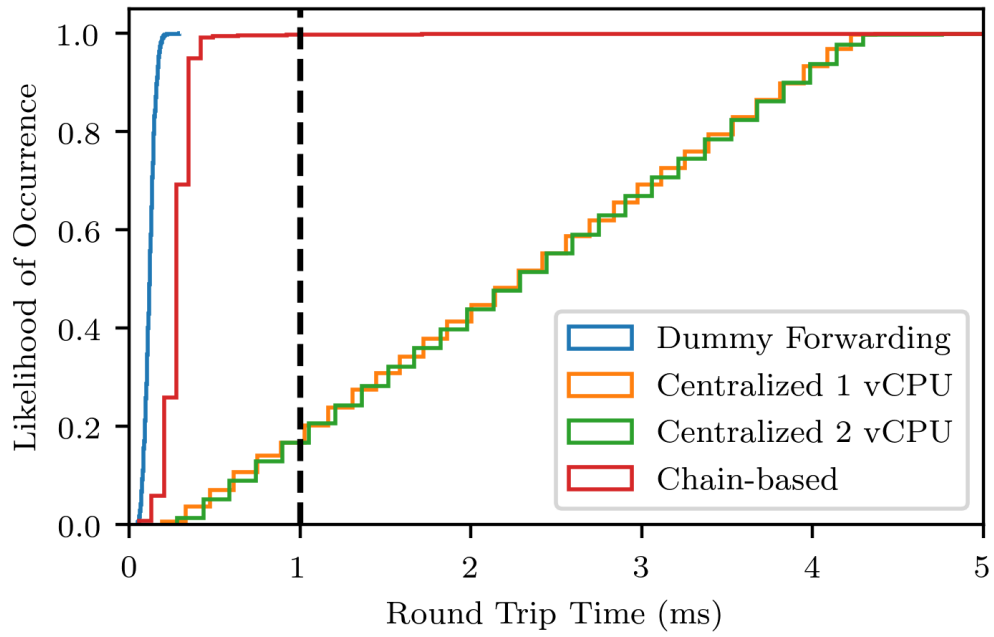


Figure 3: Service Loop Latency

## 2.1 Centralized Approach (Try to put all VNFs in a single VM)

The components of the centralized approach is showed in the Fig 4. The consideration here is to pack every VNFs in a single VM to avoid the latency introduced by transmission between VMs (We can indicate in the related work, this latency is negligible now).

[In order to make multiple VNFs cooperate properly], both kernel and user space tools need to be used. As illustrated in the Fig, the packet needs to be transmitted between kernel space and user space at least 4 times (red lines mean slow path or bottleneck). Additional copies of packets between kernel and user space introduce non-negligible latency especially in a virtual guest operating system. Besides this, the resources, especially virtual CPU, need to be shared between multiple processes both in kernel and user space. These context switching costs also time. Furthermore, the cache behavior of the CPU can not be optimized because of this context switching. Some processing related instructions can not be fetched and consistently stored in the CPU cache.




Figure 4: Centralized Approach

**Pros**

- The VNF applications can be programmed with common socket interface, so some legacy applications can be deployed without modification. [This is the main benefit, low complexity].
- Can utilize scheduler provided by the guest operating system.
- Can utilize features provided from kernel and user space tools at the same time.

### Cons

- High latency.
- Not scalable in terms of latency. As show in the RTT results. Even if two virtual CPU are used (one core handle receiving packets and another core handle sending packets out), the latency can not be reduced.
- Not flexible for dynamically resource allocation. Since several VNFs will run the same VM, the VM should be provisioned with redundant resources to ensure QoS if more VNFs are added in the chain.

## 2.2 Chain-Based Approach

The in this paper proposed chain-based approach is illustrated in the Fig 5. The basic idea here is to fully separate kernel and user space for each VNF. So each VNF should run either purely in kernel space or in the user space to avoid the latency overhead introduced by centralized approach. So each VNF runs on a dedicated VM with its own resources. These VNFs are then connected by high-performance virtual bridges with negligible transmission latency.

### Pros

- Significant reduce End-to-End latency.
- Flexible: New VNFs can added by extending current chain. Running VNFs can be removed or replaced dynamically.
- Full control of each VNF, each VNF has its own dedicated resource and runs purely in kernel or user space. Then each VNF can be optimized separately for its running environment for low latency without impact on other VNFs.

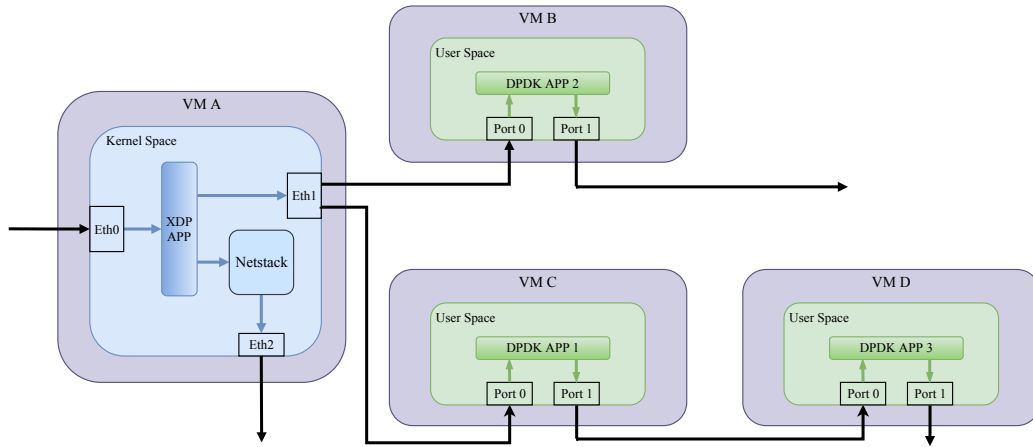


Figure 5: Chain-Based Approach

## Cons

- Uses more resources. More VMs need to be created.
- The legacy applications need to be modified to be able to run purely in the user space.
- The scheduling and resource management need to be handled by the SFC orchestrator itself instead of the operating system.
- Reduce the bandwidth. Show in the Fig 6. Because the application is fully optimized for low latency, some in kernel used strategies e.g. batching processing, are not used at all.




Figure 6: Bandwidth comparison among approaches